

CREDIT CARD FRAUD DETECTION SYSTEM

GROUP-19

GROUP MEMBERS:

Advaith Dev Krishnan -AM.EN.U4CSE20003

Divya Bisht -AM.EN.U4CSE20021

Hari Chandan -AM.EN.U4CSE20034

Likhith Reddy -AM.EN.U4CSE20039

Rithwik Sai -AM.EN.U4CSE20057

Importing the Dependencies

In [3]:

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

In [4]:

```
# Loading the dataset to a pandas dataframe
credit_card_data = pd.read_csv("C:/Users/haric/(S-5)/-Courses/data science/DS project/creditcard.csv")
```

In [5]:

```
# first five rows of the dataset
credit_card_data.head()
```

Out[5]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.

5 rows × 31 columns



In [6]:

```
#last five rows of the dataset
credit_card_data.tail()
```

Out[6]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.111864	1.014480	-0.50
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	...	0.214205	0.924384	0.012463	-1.01
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	...	0.232045	0.578229	-0.037501	0.64
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	...	0.265245	0.800049	-0.163298	0.12
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	...	0.261057	0.643078	0.376777	0.00

5 rows × 31 columns



In [7]:

```
# some info about the data
credit_card_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Time        284807 non-null float64
 1   V1          284807 non-null float64
 2   V2          284807 non-null float64
 3   V3          284807 non-null float64
 4   V4          284807 non-null float64
 5   V5          284807 non-null float64
 6   V6          284807 non-null float64
 7   V7          284807 non-null float64
 8   V8          284807 non-null float64
 9   V9          284807 non-null float64
10  V10         284807 non-null float64
11  V11         284807 non-null float64
12  V12         284807 non-null float64
13  V13         284807 non-null float64
14  V14         284807 non-null float64
15  V15         284807 non-null float64
16  V16         284807 non-null float64
17  V17         284807 non-null float64
18  V18         284807 non-null float64
19  V19         284807 non-null float64
20  V20         284807 non-null float64
21  V21         284807 non-null float64
22  V22         284807 non-null float64
23  V23         284807 non-null float64
24  V24         284807 non-null float64
25  V25         284807 non-null float64
26  V26         284807 non-null float64
27  V27         284807 non-null float64
28  V28         284807 non-null float64
29  Amount      284807 non-null float64
30  Class       284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

In [8]:

```
#checking the number of missing values in each column
credit_card_data.isnull().sum()
```

Out[8]:

```
Time      0
V1         0
V2         0
V3         0
V4         0
V5         0
V6         0
V7         0
V8         0
V9         0
V10        0
V11        0
V12        0
V13        0
V14        0
V15        0
V16        0
V17        0
V18        0
V19        0
V20        0
V21        0
V22        0
V23        0
V24        0
V25        0
V26        0
V27        0
V28        0
Amount     0
Class      0
dtype: int64
```

In [9]:

```
# Distribution of Legit transaction and fraudulent transaction
credit_card_data['Class'].value_counts()
```

Out[9]:

```
0    284315
1      492
Name: Class, dtype: int64
```

Dataset is highly unbalanced

0--> Normal Transaction

1-->Fraudulent Transaction

In [10]:

```
#separating the data for analysis
legit = credit_card_data[credit_card_data.Class == 0]
fraud = credit_card_data[credit_card_data.Class == 1]
```

In [11]:

```
legit.shape
```

Out[11]:

```
(284315, 31)
```

In [12]:


```
fraud.shape
```

Out[12]:

```
(492, 31)
```

In [13]:

```
import matplotlib.pyplot as plt
import seaborn as sns
fig, ax=plt.subplots(figsize=(6, 4))
ax=sns.countplot(x='Class', data=credit_card_data)
plt.tight_layout()
```



In [14]:

```
#statistical measures of the data for each type of class.
legit.Amount.describe()
```

Out[14]:

```
count    284315.000000
mean       88.291022
std       250.105092
min         0.000000
25%         5.650000
50%        22.000000
75%        77.050000
max      25691.160000
Name: Amount, dtype: float64
```

In [15]:

```
fraud.Amount.describe()
```

Out[15]:

```
count    492.000000
mean     122.211321
std      256.683288
min       0.000000
25%       1.000000
50%       9.250000
75%      105.890000
max     2125.870000
Name: Amount, dtype: float64
```

In [16]:

```
# compare the values for both transaction
credit_card_data.groupby('Class').mean()
```

Out[16]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9 ...	V20	V21	V22
Class													
0	94838.202258	0.008258	-0.006271	0.012171	-0.007860	0.005453	0.002419	0.009637	-0.000987	0.004467 ...	-0.000644	-0.001235	-0.000024
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.397737	-5.568731	0.570636	-2.581123 ...	0.372319	0.713588	0.014049

2 rows × 30 columns

Under-Sampling to resolve the problem of unbalanced dataset

Build a sample dataset containing similar distribution of normal transaction and Fraudulent Transaction

we would randomly take 492 data from legit dataset and we would take the fraud dataset completely and then use them to train the model

Number of Fraudulent Transaction = 492

In [17]:

```
legit_sample = legit.sample(n=492)
```

Concatenating two Dataframes using 'concat' function.

In [18]:

```
new_dataset = pd.concat([legit_sample, fraud], axis=0)
```

In [19]:

```
# first five rows of the dataset
new_dataset.head()
```

Out[19]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9 ...	V21	V22	V23
173202	121384.0	1.959801	0.016050	-1.390979	1.156715	0.372663	-0.522305	0.234468	-0.225080	0.272641 ...	0.287826	0.845287	-0.016184
7619	10540.0	1.320260	0.479284	0.110144	0.697217	0.146416	-0.544057	0.005614	-0.257043	1.309131 ...	-0.488651	-1.133276	0.032352
279855	169135.0	-2.475069	1.708159	-0.362032	-1.384144	0.475650	0.517683	0.365589	-2.176044	-0.338360 ...	2.137680	-1.177631	0.069729
271985	164855.0	2.013716	0.439072	-2.498592	1.265654	1.248825	-0.390429	0.541139	-0.141226	-0.057473 ...	-0.018697	0.099682	-0.056729
72054	54548.0	-1.130235	-0.396668	1.326885	-1.521651	-0.239949	-1.149648	0.160416	0.166297	-1.204608 ...	-0.034319	-0.122656	0.113664

5 rows × 31 columns

In [20]:

```
#Last five rows of the dataset
new_dataset.tail()
```

Out[20]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	
279863	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	-0.882850	0.697211	-2.064945	...	0.778584	-0.319189	0.639419	-0.29
280143	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	-1.413170	0.248525	-1.127396	...	0.370612	0.028234	-0.145640	-0.08
280149	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	-2.234739	1.210158	-0.652250	...	0.751826	0.834108	0.190944	0.03
281144	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	-2.208002	1.058733	-1.632333	...	0.583276	-0.269209	-0.456108	-0.18
281674	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	0.223050	-0.068384	0.577829	...	-0.164350	-0.295135	-0.072173	-0.45

5 rows × 31 columns

Checking whether class types are equal or not.

In [21]:

```
new_dataset['Class'].value_counts()
```

Out[21]:

0 492
1 492
Name: Class, dtype: int64

In [22]:

```
new_dataset.groupby('Class').mean()
```

Out[22]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V20	V21	V22	
Class															
0	92367.105691	0.076783	0.082529	0.149726	0.099438	-0.046944	-0.015630	0.078347	-0.014006	0.103854	...	-0.012634	-0.012197	-0.021440	0
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.397737	-5.568731	0.570636	-2.581123	...	0.372319	0.713588	0.014049	-0

2 rows × 30 columns

In [23]:

```
# Correlation
# Correlation is useful to find peers of input field so we are aware when building models, either to transform them (principal component)

# In this particular data, it is told that all the V-variables are already principal components. So correlation is not useful to inspect t

# But one thing we can do is to find some correlation between V-variables and Class. Recall now that coding Class into a integer data type

# code for correlation

new_dataset.corr()

# Splitting the data into features and target variable
X = new_dataset.drop(columns='Class',axis=1)
Y = new_dataset['Class']

print(X)
```

	Time	V1	V2	V3	V4	V5	V6	\
173202	121384.0	1.959801	0.016050	-1.390979	1.156715	0.372663	-0.522305	
7619	10540.0	1.320260	0.479284	0.110144	0.697217	0.146416	-0.544057	
279855	169135.0	-2.475069	1.708159	-0.362032	-1.384144	0.475650	0.517683	
271985	164855.0	2.013716	0.439072	-2.498592	1.265654	1.248825	-0.390429	
72054	54548.0	-1.130235	-0.396668	1.326885	-1.521651	-0.239949	-1.149648	
...	
279863	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	
280143	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	
280149	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	
281144	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	
281674	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	
	V7	V8	V9	...	V20	V21	V22	\
173202	0.234468	-0.225080	0.272641	...	-0.157721	0.287826	0.845287	
7619	0.005614	-0.257043	1.309131	...	-0.071264	-0.488651	-1.133276	
279855	0.365589	-2.176044	-0.338360	...	-1.057272	2.137680	-1.177631	
271985	0.541139	-0.141226	-0.057473	...	-0.215777	-0.018697	0.099682	
72054	0.160416	0.166297	-1.204608	...	-0.384260	-0.034319	-0.122656	
...	
279863	-0.882850	0.697211	-2.064945	...	1.252967	0.778584	-0.319189	
280143	-1.413170	0.248525	-1.127396	...	0.226138	0.370612	0.028234	
280149	-2.234739	1.210158	-0.652250	...	0.247968	0.751826	0.834108	
281144	-2.208002	1.058733	-1.632333	...	0.306271	0.583276	-0.269209	
281674	0.223050	-0.068384	0.577829	...	-0.017652	-0.164350	-0.295135	
	V23	V24	V25	V26	V27	V28	Amount	
173202	-0.016184	0.612282	0.310599	-0.447180	-0.003685	-0.038164	44.80	
7619	0.032352	-0.557525	0.319637	0.107080	-0.047788	0.016557	0.89	
279855	0.069729	0.172081	0.979678	0.015022	-0.442540	-0.051430	79.00	
271985	-0.056729	0.088720	0.498872	-0.511308	-0.008039	-0.037120	10.03	
72054	0.113664	0.361958	-0.106352	0.296521	0.005758	0.147132	100.90	
...	
279863	0.639419	-0.294885	0.537503	0.788395	0.292680	0.147968	390.00	
280143	-0.145640	-0.081049	0.521875	0.739467	0.389152	0.186637	0.76	
280149	0.190944	0.032070	-0.739695	0.471111	0.385107	0.194361	77.89	
281144	-0.456108	-0.183659	-0.328168	0.606116	0.884876	-0.253700	245.00	
281674	-0.072173	-0.450261	0.313267	-0.289617	0.002988	-0.015309	42.53	


[984 rows x 30 columns]

Type *Markdown* and LaTeX: α^2

We used the matplotlib and seaborn libraries to create a plot visualizing the correlation between the features of a dataset.

In [24]:


```
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(20,20))
plt.title('Credit Card Transactions features correlation plot (Pearson)')
corr = new_dataset.corr()
sns.heatmap(corr,xticklabels=corr.columns,yticklabels=corr.columns,linewidths=.1,cmap="YlGnBu")
plt.show()
```



Splitting the data into features and targets

In [25]:


```
# Histogram for all the variables  
new_dataset.hist(figsize = (20, 20))  
plt.show()
```



In [26]:

```
# box plot for all the variables
new_dataset.plot(kind='box', subplots=True, layout=(8,4), sharex=False, sharey=False, figsize=(20,20), title='Box Plot for each input variable')
plt.savefig('creditcard_box')
plt.show()
```

Box Plot for each input variable



In [27]:

```
# Prepare train, validation, test dataset

# Splitting the data into training and testing data

X = new_dataset.drop(columns = 'Class',axis =1)
Y = new_dataset['Class']
```

In [28]:

```
## normalize numeric variables
## We have divided the dataset into 80% training data and 20% testing data.
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(X)
X = scaler.transform(X)

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=2)

print('X_train shape: ', X_train.shape)
print('X_test shape: ', X_test.shape)
print('Y_train shape: ', Y_train.shape)
print('Y_test shape: ', Y_test.shape)
```

```
X_train shape: (787, 30)
X_test shape: (197, 30)
Y_train shape: (787,)
Y_test shape: (197,)
```

In [29]:

```
print(X)

[[ 0.73270511  0.77969632 -0.50369123 ... -0.09310331 -0.17261576
  -0.20814817]
 [-1.59927477  0.66393041 -0.37668314 ... -0.13661072 -0.04123192
  -0.37403617]
 [ 1.73730952 -0.02307683 -0.03975405 ... -0.52603272 -0.20446777
  -0.07894368]
 ...
 [ 1.74185381  0.30255384 -0.19926824 ...  0.29043926  0.38566857
  -0.08313716]
 [ 1.75479242 -0.13870174 -0.34746133 ...  0.78345913 -0.69011016
   0.54818926]
 [ 1.76282909  0.78552043 -0.46464129 ... -0.08652096 -0.11774154
  -0.21672403]]
```

In [30]:

```
print(Y)

173202    0
7619      0
279855    0
271985    0
72054      0
..
279863    1
280143    1
280149    1
281144    1
281674    1
Name: Class, Length: 984, dtype: int64
```

Splitting the data into training data and test data

In [31]:

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.2,stratify = Y, random_state=3)
```

In [32]:

```
print(X.shape,X_train.shape,X_test.shape)

(984, 30) (787, 30) (197, 30)
```

```
## Model Training

1) LogisticRegression
```

In [33]:

```
model=LogisticRegression()
```

In [34]:

```
# training the Logistic Regression model with the training data
model.fit(X_train,Y_train)
```

Out[34]:

```
LogisticRegression()
```

In [35]:

```
#accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction,Y_train)
```

In [36]:

```
print('Accuracy on Training data : ',training_data_accuracy*100)
```

Accuracy on Training data : 95.1715374841169

In [37]:

```
#accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction,Y_test)
```

In [38]:

```
print('Accuracy on Test data : ',test_data_accuracy*100)
```

Accuracy on Test data : 94.9238578680203

In [39]:

```
# KNN model

from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=7)
knn.fit(X_train,Y_train)
knn.score(X_test,Y_test)

# accuracy on training data

X_train_prediction = knn.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction,Y_train)

print('Accuracy on Training data : ',training_data_accuracy*100)

#accuracy on test data

X_test_prediction = knn.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction,Y_test)

print('Accuracy on Test data : ',test_data_accuracy*100)
```

Accuracy on Training data : 93.01143583227446

Accuracy on Test data : 91.87817258883248

In [40]:

```
# SVM model

from sklearn import svm
svm_model = svm.SVC(kernel='linear')

# training the Support Vector Machine model with the training data
svm_model.fit(X_train,Y_train)

#accuracy on training data
X_train_prediction = svm_model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction,Y_train)

print('Accuracy on Training data : ',training_data_accuracy*100)

#accuracy on test data
X_test_prediction = svm_model.predict(X_test)

test_data_accuracy = accuracy_score(X_test_prediction,Y_test)

print('Accuracy on Test data : ',test_data_accuracy*100)
```

Accuracy on Training data : 95.04447268106735

Accuracy on Test data : 95.43147208121827